



Linköping University

DEPARTMENT OF COMPUTER AND INFORMATION SCIENCE

TDTS06 COMPUTER NETWORKS

Distance Vector Routing

Report

Group C/D 13

CHVÁTAL Martin march011@student.liu.se

PESCHKE Lena lenpe782@student.liu.se

Teaching assistant :

SCHMIDT Johannes johannes.schmidt@liu.se

October 12, 2014

I How distance vector routing works

In general Distance vector routing uses a decentralized routing algorithm to find the least-cost path between routers in a network. The distance vector routing algorithm considers a network as an undirected graph of routers, represented as nodes, and links, represented as edges. It is iterative, asynchronous and distributed [KR13, p. 371]. It makes use of the Bellman-Ford equation:

$$d_x(y) = \min_v \{c(x, v) + d_v(y)\} \quad (1)$$

The equation states that the cost of the least-cost path from node x to y , $d_x(y)$, is the minimum of the sum of the cost to get from x to a neighbor v and the cost of the least-cost path from v to y .

The algorithm works by making the nodes exchange distance vectors. To make use of the equation, each node needs to know about all of the other nodes in the network and the cost to its direct neighbors. Let's walk through the algorithm step by step.

1. Initially, each node knows the cost to its direct neighbors. It sends its distance vector to its neighbors.
2. Upon receiving an update (link cost change or new distance vector from a neighbor), a node recomputes its distance vector using the Bellman-Ford equation (1). If there is a change, it sends its own new vector to the all the neighbors.
3. "The process of receiving updated distance vectors from neighbors, recomputing routing table entries, and informing neighbors of changed costs of the least-cost path to a destination continues until no update messages are sent." [KR13, p. 375]

After step 3, the estimated costs have converged to the actual least-costs.

II How we tested the algorithm

wait for the code

III Poisoned reverse

When a link cost in the network increases, a count-to-infinity problem can occur. A packet can then be caught in a routing loop until the nodes have figured out the shortest path, which can take a while (as many iterations as the difference between the previous and the new least-cost). Poisoned reverse is an addition to the DVR algorithm which helps avoiding this situation. If a node x routes through y in order to get to z , x lies to y about its distance to z , so that an increase in the cost link between y and z results in a convergent update that takes only two iterations.

Where it fails Poisonous reverse fails if a count-to-infinity problem triggers a loop of three or more nodes [KR13, p. 378].

Let us consider a network of 4 routers, A , B , C and D . The network is represented below in Figure 1.




Figure 1: Example topology where poisoned reverse fails.

	<i>via</i>		
	B	C	D
B	1	2	∞
C	2	1	∞
D	2	3	∞

Table 1: A's initial routing table

	A	C	D
A	1	2	∞
C	2	1	∞
D	∞	∞	1

Table 2: B's initial routing table

	A	C	D
A	1	2	∞
C	2	1	∞
D	∞	∞	50

Table 3: B's routing table after step 1

	<i>via</i>		
	B	C	D
B	1	2	∞
C	2	1	∞
D	51	3	∞

Table 4: A's routing table after step 2

Solutions In practise, the Routing Information Protocol (RIP) uses DVR with poisoned reverse where infinity is 16 hops. This does not avoid the count-to-infinity problem, but limits the damage by stopping updates after it reaches ‘infinity’. In addition, *hold-down timers* can be used: if a route becomes unusable, it is kept in the tables with an infinity value for a certain amount of time. If an update advertises

	A	C	D
A	1	2	∞
C	2	1	∞
D	4	∞	50

Table 5: B's routing table after step 3

a better or equally good route, it is replaced. On the other hand, if the update indicates a worse cost, it is ignored.

Another solution (which actually solves the problem) is using *path vectors*: in addition to the distance vectors, we can keep a record of the path, i. e. the nodes traversed to obtain the cost. When a link cost increases or a link disappears from the network, a node can check if the path is still valid by looking it up. This is used in BGP [Siu99].

References

- [KR13] James F. Kurose and Keith W. Ross. *Computer Networking: A Top-Down Approach*. Pearson, 6th edition, 2013.
- [Siu99] Prof. K.-Y. (Sunny) Siu. Distance-vector routing (distributed Bellman-Ford algorithm). <http://web.mit.edu/course/2/2.993/www/lectures/lecture14.txt>, April 1999.